

Fine-grained role-based access control by encryption

Andreas Severin Hauch Trøstrup
Computer Science
IT University of Copenhagen
atro@itu.dk

Lucas Frey Torres Hanson
Computer Science
IT University of Copenhagen
luha@itu.dk

Abstract—Data lakes supporting multiple query engines lack fine-grained access control; only “fully managed all-in-one cloud platforms” like Snowflake or Databricks offer it, within their own ecosystem.

We propose Object Store Wrapper Service (OSWS), a system enforcing column-level role-based access control at the Object Store layer using Parquet Modular Encryption with envelope encryption, and still supporting query engines that are already S3 compatible.

Our evaluation shows that OSWS adds an acceptable overhead for tiny to small Parquet files, when using a DEK cache, with warm p95 GET latency of 30ms for a 5MB Parquet file, but with larger files it becomes unusable, and design choice changes would allow for an efficient solution.

OSWS demonstrates that encryption-based access control at the Object Store layer is a viable approach for data lakes, though limitations exist for query engines that cache Parquet metadata; however, those issues are addressable.

TODO LIST

- ✓ In **methodology**, add info about setup for testing DuckLake and what the issues were
- System Design has to be refactored to not mention code files
- ✓ In **Discussion**, we have to add more info about why Parquet Sharp is not the way forward, as it does most work sequentially, and our other idea would allow for pre-fetching the keys and doing work asynchronously
- ✓ Why is footer encryption not enabled.
- ✓ Fix the not-defined times for how slow it is
- ✓ Lucas uses p99 for micro stuff. Use p95 instead to match e2e
- Include figure for parquet with unencrypted footer in appendix
- ✓ LRU has to be mentioned in system design
- ? Seems like there is a lot of repetition in the beginning of Section VII
- ✓ Have to mention AES-CTR in the discussion as it is currently only referred to in the conclusion.

I. INTRODUCTION

Today, when you want to use a data lake, you have two options. *Option 1*: Select a “fully managed all-in-one cloud platform” that integrates with data lakes and provides fine-grained access control. *Option 2*: Provide the external query engines the ability to interact with the data lake, but lose the

ability for fine-grained access control. So either you choose *Option 2* and limit who the data can be shared with, based on whether they should only be able to access part of the data, or choose *Option 1* and limit the query engines the data scientists want to use.

Object Stores, such as Amazon S3, which can be used as data lakes, do not understand a granularity finer than file level. This limitation prevents external access control services from governing access to a finer granularity, such as column-based. If, instead, the mechanism for access control was moved within the data lake itself, this could enable fine-grained access control. In the research project “Reviewing options for fine-grained Role-Based Access Control in Data Lakes”, Trøstrup and Lucas[1], the solution called Object Store Wrapper Service (OSWS) was proposed.

The idea of OSWS is to encrypt the Parquet files stored inside the data lake using Parquet Modular Encryption (PME), which supports encrypting each column. The keys will be stored in a key vault and in wrapped format within the Parquet files, and will be mapped to roles, to limit which roles have access to specific columns. OSWS will expose an S3-compatible API, which can then decrypt the data for the “fully managed all-in-one cloud platforms” and the external query engines when getting normal S3 requests. Thus, regardless of where to read the data from, it will have been ensured that they can only see what they are supposed to.

In this thesis, a PoC of OSWS will be created to prove that a system supporting third-party query engines, such as DuckDB, without modifications, can use OSWS, though, as quickly realized, interoperability with the “fully managed all-in-one cloud platforms” is partly a job that the respective companies have to allow. After developing the PoC and coupling a query engine to it, experiments will be conducted to evaluate the performance and **security of the system**, compared to available solutions. Metrics include end-to-end read and write latency, as well as the cryptographic overhead, in relation to the size of the Parquet file, introduced by the OSWS. Experiments will use data sizes representative of realistic data sizes.

A. Research Questions

Is it possible to create a wrapper service that ensures access control on column-level using RBAC and has interoperability with existing query engines and “fully managed all-in-one cloud platforms” without modifying them?

B. Objectives

This research paper will have the following objectives:

- 1) Create OSWS PoC using an underlying S3-compatible Object Store
- 2) Ensure OSWS enforces access control using RBAC outward using encryption
- 3) Make it compatible with “fully managed all-in-one cloud platforms”
- 4) Make it compatible with query engines
- 5) Benchmark OSWS to measure the latency it adds.
- 6) Reflect on design decisions and which alternatives would have been better

II. RELATED WORK

A. Reviewing options for fine-grained Role-Based Access Control in Data Lakes

This thesis is a continuation of the research paper titled “Reviewing options for fine-grained Role-Based Access Control in Data Lakes”[1].

This showed that there is a gap within the data lake space. One option is to use “fully managed all-in-one cloud platforms”, such as Databricks Unity Catalog or Snowflake, to provide granular control of what clients can see, with the platform’s own built-in query engine. The other option is to use solutions using vended credentials, such as Lakekeeper or Apache Polaris; essentially, providing temporary credentials to clients, who use whatever external query engine they want to use, but in exchange, only file-level granularity can be controlled.

Based on this, OSWS was proposed to be able to provide both columnar access control and let clients use whatever query engine they wanted to use.

This was the foundational work to reflect on how OSWS could theoretically work. After beginning to implement OSWS, discoveries revealed that it did not work as initially planned, and as described in the research paper.

The biggest change from the research paper is that OSWS does not support external query engines decrypting locally, but has to get the decrypted Parquet through OSWS. As also mentioned later, “fully managed all-in-one cloud platforms”, such as Snowflake or Databricks Unity Catalog, are not guaranteed to be supported as they would have to whitelist a URL on which OSWS runs, and have thus not been able to be tested. Lastly, data lake-specific systems, such as DuckLake, are not supported either, due to it caching metadata and the Parquet file being modified in OSWS, discussed in more depth in Section VII. Though this is an issue that could be removed in the future, it will require OSWS to have another system design than the one it currently has.

B. Membrane: A Cryptographic Access Control System for Data Lakes

Membrane, referenced in Trøstrup and Lucas[1], is a proposal to ensure FGAC at the cell level and encryption-at-rest in data lakes using a custom format, Kumar et al.[2], thus resembling what OSWS sets out to do. Besides making the

original client responsible for handling the original encryption keys to provide access to others, it also forces the query engine to **download the whole file** to decrypt it with the provided view on its machine. The advantage of this is that once the query engine has the data, given the view, they only have to make computations on the relevant data. This differs from OSWS, where it actually runs computations on all data, but instead only gives the query engine the finalized data. The computation volume for Membrane is therefore correlated to the size of the view size, whereas for OSWS, it is related to the file size, but OSWS supports existing query engines which can interact with S3. Membrane also has finer-granularity access-control than OSWS, cell-level vs column level, but this comes at the cost of using a proprietary format.

C. One Stone, Three Birds: Finer-Grained Encryption with Apache Parquet @ Large Scale

Shang et al.[3] present a system for fine-grained encryption of Parquet files at Uber’s scale, solving encryption-at-rest and access control as a combined problem—similar to OSWS. Their approach uses column-specific encryption keys stored in KMS, with readers and writers performing cryptographic operations directly through a Parquet library extension.

The key difference is where the cryptographic operations occur. The Uber system requires extending the Parquet reader/writer libraries used by query engines, making it intrusive to adopt. OSWS performs all cryptographic operations in the Encryption Gateway, exposing a standard S3 API so that unmodified query engines can access data without library changes. The trade-off is that OSWS introduces a network hop and full-file processing overhead, while Uber’s approach avoids both but requires per-engine integration.

III. BACKGROUND

Data lakes are not a new concept, but the term was first coined in 2011 by Dixon[4], and some of the terminology is a bit inconsistent. This section will describe the terminology used within this thesis to ensure consistency. A lot of the definitions will be taken directly from the previously written research paper, which will be referenced and quoted accordingly.

A. OSWS

Object Store Wrapper Service (OSWS) is the concept of a layer on top of an Object Store and which support S3-compatible API endpoints. It is meant to provide fine-grained role-based access control by encryption at the columnar level. This will be achieved by using Parquet Modular Encryption (PME), and will only decrypt the columns which are authorized to the client requesting a given Parquet file.

B. JIT Provisioning

Just-In-Time (JIT) Provisioning is an identity management process handling creating user accounts when they are needed. [5] This is needed during the initial user creation, see Section IV.C.g.

C. RBAC

In Trøstrup and Lucas[1], Role-Based Access Control (RBAC) was defined as follows:

Role-based Access Control (RBAC) [6] is an approach to access control based around assigning permissions to roles, and grouping users into these roles. Access control is then managed by controlling what roles are allowed to do what, instead of managing each user individually. The “Core RBAC” model, as defined by Ferraiolo et. al.[7], consists of five basic elements: *users*, *roles*, *permissions*, *operations* and *objects*. A *user* is some actor, typically a human user, but could also be an autonomous actor, who needs to be authorized for some set of actions. A user is assigned one or more *roles*. A role is simply some named collection of authorizations; for example, a job title. A role is assigned one or more *permissions*, which are what give authorization to perform an *operation* on a secured *object*. [...] In addition, Core RBAC also describes *sessions*. A *session* is a mapping of one user to a subset of the roles they are authorized for; i.e. a user “activates” one or more roles that have the necessary permissions, when they need to perform a specific operation on an object. This allows for finer control of roles that are active in a given context, to prevent unnecessary access levels. Commonly, a hierarchy is introduced to allow for structuring roles which should inherit some base permissions; for example, a Doctor and Nurse role might inherit permissions from some base “Healthcare Worker” role. This is referred to as “Hierarchical RBAC” [7].

— Trøstrup and Lucas[1]

D. OIDC

OpenID Connect (OIDC) is an identity authentication protocol based on the authorization OAuth 2.0 framework.[8] It provides developers with a standardized, simple way to verify the identity of users trying to access web applications. It allows users to authenticate using their existing accounts previously registered with an OpenID Provider, such as Microsoft Entra ID[9] or PocketID[10], which the web application can then contact to verify their identity. Thus, it eliminates the need to implement an authentication layer in the web application.

E. KMS/KV

Key Management Service (KMS) and Key Vault (KV) are both services used to manage keys.[11], [12] KMS is widely used within data storage, and is what AWS use to refer to their system, whereas Azure uses KV. This paper will use the term KV, as OSWS has been set up against Azure KV but could have used KMS. KV is a way to store cryptographic keys. When the cryptographic keys are within the KV, they can no longer be retrieved, and thus if data has to be encrypted/decrypted, it either has to be sent to the KV or envelope encryption can be used, so the KV is responsible for unwrapping the DEK, see more in Section III.I.

F. Data Lake

In the previously written research paper, data lakes were defined as follows:

A “data lake” is a centralized repository of data, which could be structured, semi-structured or unstructured.[4] It differs from a traditional database in that a data lake separates compute and storage layers. The data itself is stored in object stores like Amazon S3[13], in file formats like CSV, Parquet¹, etc. The compute nodes are simply any query engine, such as Apache Spark [14], that connects to the object store and uses the data. This could be a data scientist working on their own machine, who connects to a data lake and runs queries; or it could be a part of a managed all-in-one platform like Snowflake, see more in Section III.K. As the object store just contains raw data, both structured and unstructured, this can make it more complicated to query, especially for large data lakes. One approach to querying is “schema on read”, in which the query engine infers a schema at query time based on the files in the data lake. This is in contrast to “schema-on-write” used in traditional relational databases, where the schema is strictly enforced when writing. This doesn’t work well for data lakes due to the nature of storing raw data in object stores. Typically, data lakes introduce a *catalogue* on top of the object store that manages file schemas, metadata and snapshots, allowing for structured reads and more performant querying through techniques like partitioning, i.e. filtering unnecessary data. This has also been referred to as a *data lakehouse* .[15]

— Trøstrup and Lucas[1]

G. Parquet

In the previously written research paper, Parquet files were defined as follows:

Parquet [16] is a columnar file format supported by many data processing systems.

A Parquet file is structured using row groups, which contain columns, along with a footer which contains metadata. The columns are structured as column chunks, which contain several *data pages*. A data page contains a header, which describes information like the size of the page, and following that, the actual values as bytes. The file footer contains metadata that specifies what is stored in the file and where, that is, where a reader should look to find a specific column, or the statistics of a column (e.g. min/max). [...] The objects that hold metadata, like statistics, and structural information, including physical offsets, such as the footer and page headers, are serialized using Apache Thrift[17], a language-agnostic serialization framework, and referred to as “Thrift Structs”. Essentially, the object that contains

¹Section III.G defines Parquet




Figure 1. Illustration of the Parquet file structure using Plaintext Footer. The red key denotes what is encrypted with the column key. From [18].

the information, for example, a `PageHeader` object [...] or a `ColumnMetaData` object [...], is serialized into bytes using Thrift, which can then later be restored to the same object.[18]) Then, to read a Parquet file, the reader de-serializes the Thrift structures to get the necessary metadata, and then reads the actual data values using the information from the metadata.

— Trøstrup and Lucas[1]

H. PME

Parquet Modular Encryption (PME) is a way to allow granular control of how the Parquet file data and metadata should be encrypted. Singular columns can be encrypted, and the footer as well. For OSWS, footer encryption is not enabled, as it was outside of the scope due to not containing sensitive column-related data. When encryption is used, serialized Thrift structs are encrypted using the given cryptographic key, and the data pages themselves are encrypted as well. In OSWS, it is saved as a tuple in the footer, containing a reference to the KEK in KV and the wrapped DEK. This allows OSWS to get the DEK unwrapped to then decrypt the column.[18] **Wrong and move to system design**

a) Encryption in Parquet:

In PME, column encryption is always enabled, but it provides two options for footer encryption; encrypted or plaintext footer. Figure 1 shows a Parquet file with plaintext footer enabled. In plaintext footer mode, a `ColumnCryptoMetaData` struct is added to the each column chunk, which contains the necessary metadata to decrypt the column, such as key ID and algorithm. This is show on

the right side of Figure 1. Notably, each column's statistics are moved to an encrypted `ColumnMetaData` struct, also shown on the right side of Figure 1. This ensures that no information about the values can be gained without decrypting the columns. Non-sensitive information like the offsets and other metadata is stored as-is without encryption.

I. Envelope Encryption (KEK, DEK)

KV does not support retrieving keys, and as a result, OSWS uses envelope encryption. This means that during encryption of a column, cryptographic keys for each are generated, called a data-encryption-key (DEK), which is used to encrypt the columns. Then, for the single Parquet file, a single key-encryption-key (KEK) is generated, resulting in less data compared to if each DEK had its own separate KEK within KV. This KEK then encrypts all the DEKs, also called wrapped DEKs, and gets stored in KV. The wrapped DEKs are then stored in the metadata for the Parquet file together with a KEK id. When the columns are to be decrypted, the wrapped DEKs are sent to KV to be unwrapped, and the columns can be decrypted.[19]

J. TTL

Time-to-live (TTL) is a way to define the lifetime of some item within a dataset. When either the item's TTL expires or the dataset is full, the item closest to its TTL is removed. [20] In OSWS, TTL defines the lifetime of unwrapped DEKs in memory. When the TTL is reached, the data will be removed and has to be re-fetched from the KV when needed again. OSWS defines two different TTLs as a way to follow general security guidelines, which can be read about in "NIST

Special Publication 800-57 Revision 5 Recommendation for Key Management”[21], where OSWS uses a TTL of five minutes for DEKs authorized to admins, and the rest have a TTL of 15 minutes.

K. Fully managed all-in-one cloud platforms

In this paper, “Fully managed all-in-one cloud platforms” will be used to define solutions such as Snowflake and Databricks. Access control within those solutions is managed by the query engines together with the catalogue. This means the query engine is trusted to filter out rows that are not permitted, and it then masks or removes those values, which is possible as the platforms themselves manage it.

L. External Query Engines

External query engines reference query engines that operate outside the data layer. This is one of the traits of data lakes, as the computations and data are split up. These can include PyArrow, PySpark, DuckDB, and DuckLake. Here, it is important, as later discussed, to distinguish that DuckLake uses schema on write. This means that this specific query engine, once it creates a file, e.g. a Parquet file, it saves the schema and metadata within its internal catalogue, whereas the other solutions use schema on read, and thus ask for the metadata during query time.

IV. SYSTEM DESIGN

What the current OSWS solves is that query engines, which read metadata from Object Store and do not use cached values, will have columnar access control enforced on the data they read. This ensures that, from the query engine’s perspective, they are using S3 directly, but in reality use OSWS, and it can enforce whatever RBAC rules it wants on the clients. So if you have a dataset with columns `name`, `birthday`, and `relational status`, and the query engine is not granted access to `birthday`, it will receive all columns, but `birthday` will have null or another default value.

Below, the design choices will be described in more depth, together with deeper technical information.

A. Architecture Overview

OSWS is set up as a wrapper for S3 and thereby provides S3-compatible API endpoints for clients. The purpose of OSWS is to provide column-level access controls to S3, whilst not modifying the S3 API. This ensures that most S3-compatible query engines can use OSWS without modification and will only read what they are supposed to.

The system is built on ASP.NET Core 10 using the minimal API pattern[22], which replaces traditional MVC with lightweight endpoint definitions. OSWS uses: PME via. Parquet Sharp, PostgreSQL for RBAC metadata, Azure KV for key management, and a configurable S3-compatible Object Store (tested with Cloudflare R2 and Digital Oceans Spaces for benchmarking).²

TABLE I
OSWS SYSTEM COMPONENTS AND THEIR RESPONSIBILITIES

Component	Responsibility
Encryption Gateway	S3-compatible API host. Handles incoming <code>GetObject</code> / <code>PutObject</code> requests, enforces AWS Signature V4 authentication, and handles encryption, decryption, and column filtering.
Parquet Solver	Encrypts and decrypts Parquet files using PME with ParquetSharp. Generates DEKs, wraps them via the Key Manager, and masks unauthorized columns.
Key Manager	Manages cryptographic keys through Azure KV (or an internal provider). Wraps/unwraps DEKs using RSA-2048 KEKs stored in the vault.
RBAC Store	PostgreSQL database storing users, roles, role assignments, role inheritance, column permissions, and S3 credentials. Evaluated per-request to determine authorized columns.
Admin Frontend	React web UI for managing roles, permissions, and users. Uses OIDC authentication and provides a SQL-like query editor for RBAC operations.

OSWS consists of the five main components, see Figure 2 or Table I.

The **admin frontend** is for RBAC, thus managing roles, permissions, and users within the OSWS system, more specifically within the **RBAC store**, which is the database for storing these relations. Here, it can be defined what columns a user can access later on and, by extension, which columns will be masked.

The **encryption gateway** is the entry point for retrieving and putting files into the actual S3. It is responsible for delegating work to the other components, so encrypting and decrypting files, partially, is possible. Whenever it receives a GET request, it will first get the file from S3 and see which columns the client is able to read, by using the **RBAC store**. Then it will use the **Parquet solver**, which in turn will handle trying to identify which DEKs are in the Parquet file. It can then go through the **key manager** to get the unwrapped DEKs based on permission level; it will check locally first, and if not stored, go to KV to get them unwrapped there, and store them for later use. After the key manager returns the unwrapped DEKs, the Parquet solver can decrypt the columns and mask out the ones it cannot decrypt. Now the encryption gateway can return the Parquet file, showing only the information that the client has access to.

When the encryption gateway gets a PUT request, it first gets the user’s base role. For testing purposes it was tied to the user’s first role, but should have been tied to their base role or the one provided through the credential. **TROLLE DO SOMSING** Then the Parquet solver is given the file and role, so it can generate new DEKs per column and a KEK for the Parquet file. The RBAC store is updated with the new columns and key references, and the file is encrypted afterwards. It also handles wrapping the keys initially so the wrapped DEKs can be put into the metadata, and KEK is given to the key manager to insert into KV.

²Codebase available at github.com/lucasfth/osws




Figure 2. Architecture of OSWS showing components, external services, and communication paths




Figure 3. Entity Relationship Diagram showing the design of the database

This allows OSWS to ensure fine-grained role-based access control in data lakes, but of course, there are more in-depth design choices that have been made to make this work, both for better and worse.

B. Design Choices

a) Database Design:

An Entity Relationship describing the database design can be seen in Figure 3. The database is mostly used for RBAC metadata; however, it also stores credentials used for AWS Signature V4 request signing, and external identity information from the OpenID Provider(s).

b) Role Hierarchy:

OSWS also supports Hierarchical RBAC through the RoleInheritance tables. This is a self-referential table in which a hierarchical relation is created through the ParentRoleId and ChildRoleId foreign keys, so that a *parent* inherits from a *child*, see RoleInheritance in Figure 3. This means that all permissions of *child* also become permissions of *parent*. To retrieve the full set of effective roles for a user, a recursive SQL query traverses the hierarchy. This can be seen in Listing 1.

```
WITH RECURSIVE effective AS (
  SELECT ra."RoleId" AS "Id"
  FROM "RoleAssignments" ra
  WHERE ra."UserId" = {userId}
  UNION
  SELECT ri."ChildRoleId"
  FROM "RoleInheritances" ri
  JOIN effective e
  ON ri."ParentRoleId" = e."Id"
)
SELECT DISTINCT r."Id", r."Name"
FROM "Roles" r
JOIN effective e ON r."Id" = e."Id"
```

Listing 1. A recursive SQL query to get the full set of effective roles given a userId.

c) Column-Level Permissions:

Each *Permission* record maps a *Role* to a *Column*. This means that a column is decrypted if and only if one of its roles grants access to the given column. This ensures fine-grained access control. If a Parquet file *P1* with columns *name*, *age*, and *civil_personal_register* exists. Then a user *U1* and a user *U2* can respectively give access to role *R1* and *R2*. *R1* provides access to only *name* and *R2* only to *age* and *civil_personal_register*. Then *U1* would see a table full of valid *names* whilst *age* and *civil_personal_register*

are null/default values, while *U2* would see *age* and *civil_personal_register* values.

C. Caching

There are two tiers to the caching system used within OSWS with the intention of minimizing latency from remote services.

a) Encrypted File Cache:

Given the configuration of OSWS, the Parquet files fetched from the object store can be cached. It uses the LRU (Least Recently Used) policy to cache the encrypted Parquet files on the local file system. They are keyed with `SHA256(bucket::key)`, see `OSWS.ParquetSolver/Helpers/EncryptedFileCache.cs`.

b) DEK Cache:

Unwrapped DEKs are cached in-memory to avoid repeated KV calls, as it is one of the major latency attributes, discussed later. The cached DEK is keyed with the KEK identifier, and in the configuration, the TTL can be defined for regular and admin users, as admin users' keys are most often more privileged. The cache enforces a maximum capacity and evicts the expired entries first, and then the oldest entries by expiration date.

c) Envelope Encryption:

To encrypt the Parquet columns, OSWS needs to use DEKs and KEKs to achieve envelope encryption[23]. This deviates from the original proposed solution in Trøstrup and Hanson[1], as KV does not support key retrieval, and the overhead of sending a Parquet column to KV each time for decryption and encryption is high. By using envelope encryption, OSWS can cache the unwrapped DEKs with a specified TTL to allow quicker decryption. This results in a solution with lesser overhead and ensures column-level access control.

d) Parquet Solver:

Inkluder en paragraf om hvordan vi storer metadata i parquet filen ift. hvordan forskellene er fra den "rigtige" spec og hvordan vi gør det. Hvis vi har overskud ville et diagram nok være nice Parquet Solver handles the cryptographic operations for the Parquet files. Important design decisions made here are that it uses Parquet Sharp. Parquet Sharp was chosen as it supports PME, but does not support partial decryption/reads of Parquet files. This results in the fact that when a Parquet file has to be decrypted, it essentially decrypts and copies over all the columns into a new Parquet file. But it does not support copying over the columns which are not supposed to be decrypted. As a result, dummy columns are created and copied over into the new Parquet file.

This is one of the design choices made that later created more problems than it solved.

e) Authentication:

Due to S3 requiring signage, the calls from the clients going to OSWS are signed as well, and as a result, OSWS cannot reuse the signature and encrypt the Parquet columns. Internal authentications were therefore needed, and OSWS then has its own signature for S3.

f) AWS Signature V4 (S3 API):

Because OSWS had to use an S3-compatible API, it also had to use the same validation as S3. A

custom `SigV4AuthenticationHandler` has been created to parse the authorization header and extract the `AccessKeyId`. The keys corresponding to the `S3Credential` record are then looked up in the internal PostgreSQL, see the `S3Credential` table in Figure 3. The signature is verified using the stored secret key. On success, a `ClaimsPrincipal` is created carrying the user's database ID, name, email, and default role.

This follows the AWS specification "Authenticating Requests (AWS Signature Version 4)"[24].

g) OIDC:

The React frontend, for handling RBAC, uses OIDC for the authentication flow. The OIDC provider chosen for OSWS is Pocket ID, as it is simple to set up, is open-source, and uses passkeys, which are more secure than standard MFA, as "common multifactor authentication methods can be intercepted or relayed"[25].

The backend still supports other OIDC providers. To use other providers, `OSWS.WebApi/appsettings.json` needs to be updated to reflect the change, and the frontend `.env` has to point to the specified authority and client ID.

On the first login, the `api/me` endpoint triggers JIT provisioning, and a new `User` record and `ExternalIdentity` record are created in the database. Subsequent logins will synchronize the OIDC provider's claims.

D. RBAC

In Section III.C, RBAC is defined to consist of: *users*, *roles*, *permissions*, *operations*, *objects* and *sessions*. OSWS implements RBAC mostly following this, but it differs in a few ways. In Figure 3, the core RBAC entities are shown: the tables `User`, `Role`, `RoleAssignment`, `RoleInheritance`, `Column`, and `Permission`. `User` and `Role` have a many-to-many relationship through `RoleAssignment`. Likewise, `Role` and `Column` have a many-to-many relationship through `Permissions`. The `Column` table can be seen as the *object* of Core RBAC, as it is currently the only secured object to which access is controlled. There is, however, no concept of *operations*, only whether or not access is granted. There is also no formal concept of *sessions* as described by Ferraiolo et al.[6]; however, when a user is authenticated through their S3 Credential, their assigned roles are loaded. In this way, it can be seen as an activation of all the users assigned roles.

a) Key Hierarchy:

The KEK sizes are specified within the KV that the admin chooses. Azure KV is used to store the KEKs, chosen due to pricing, but is able to be switched out, and RSA-2048 is used for the keys. The KEKs are created in OSWS, but after that, the KEK never leaves KV again, and are then called by OSWS to wrap and unwrap the DEKs, which themselves are stored within the Parquet metadata. The DEKs are created within OSWS as well, use AES, and, based on the setup, can have sizes 128, 192, or 256, and are specifically created during the encryption flow.

b) Administrative Endpoints:

Endpoints for admin endpoints are restricted to only users who have the `IsRbacAdmin` flag, seen in the `User` table in Figure 3. This field is populated from the `ClaimsPrincipal` when provisioning the user in the `/api/me` route. Thus, it requires the OIDC Provider to include this claim on users who should have admin access. Alternatively, it can be set manually.

TABLE II
ADMIN ROLE MANAGEMENT ENDPOINTS PROVIDED BY OSWS

Method	Path	Description
GET	<code>/api/admin/columns</code>	Get all columns
POST	<code>/api/admin/columns/{columnId}/roles/{roleId}</code>	Grant access on column with ID <code>columnId</code> to role with id <code>roleId</code>
DELETE	<code>/api/admin/columns/{columnId}/roles/{roleId}</code>	Revoke access on column with ID <code>columnId</code> to role with ID <code>roleId</code>
GET	<code>/api/admin/roles</code>	Get all roles
POST	<code>/api/admin/roles</code>	Create a role
DELETE	<code>/api/admin/roles/{id}</code>	Delete role with ID <code>id</code>
POST	<code>/api/admin/roles/{parentId}/inherit/{childId}</code>	Make role with ID <code>parentId</code> inherit role with ID <code>childId</code>
DELETE	<code>/api/admin/roles/{parentId}/inherit/{childId}</code>	Delete inheritance of role <code>parentId</code> from role <code>childId</code>
GET	<code>/api/admin/users</code>	Get all users
POST	<code>/api/admin/users/{userId}/roles/{roleId}</code>	Grant role with ID <code>roleId</code> to user with ID <code>userId</code>
DELETE	<code>/api/admin/users/{userId}/roles/{roleId}</code>	Revoke role with ID <code>roleId</code> from user with ID <code>userId</code>

c) Frontend Architecture:

A frontend for interacting with the endpoints described in Table IV and Table II was built using Typescript-React. A user can log in using Pocket ID, as described in Section IV.C.g. Here, the user can create credentials to be used for the S3-Compatible API endpoint. If the user is an RBAC Admin, they get access to the admin panel for managing roles. The admin panel includes a table view for viewing currently existing users, roles, columns and permissions. To interact with the endpoints described in Table II, a “query editor” was built to manage RBAC operations in an SQL-like syntax, reminiscent of the syntax used to manage permissions in e.g. PostgreSQL.³ Using the JavaScript library `peggyjs` [26], a simple grammar was written to parse statements into API calls. Some example statements to configure RBAC permissions and the API calls they parse can be seen in Listing 2.

³<https://www.postgresql.org/docs/current/sql-grant.html>

```
CREATE ROLE admin;
=> POST /api/admin/roles { name:
  "admin" }
CREATE ROLE intern;
=> POST /api/admin/roles { name:
  "intern" }
GRANT intern TO ROLE admin;
=> POST /api/admin/roles/1/inherit/2
GRANT ACCESS ON name TO intern;
=> POST /api/admin/columns/1/roles/1
GRANT ACCESS ON ssn TO admin;
=> POST /api/admin/columns/2/roles/2
GRANT admin TO USER alice;
=> POST /api/admin/users/1/roles/1
```

Listing 2. Example statements and their parsed API calls

d) Encryption Flow:

Make into seq. diagram and make it more abstract and not code specific to give an overview of how OSWS handles it. Reference the diagram. Use pintora.

The encryption flow works as follows:

- 1) Client uploads unencrypted Parquet file, via PUT `/s3/{bucket}/{key}`.
- 2) OSWS creates an RSA-2048 key and is tagged with the uploading user’s role.
- 3) For each column designated for encryption, an AES DEK of specified sizes is generated and encrypted, within `OSWS.ParquetSolver/Helpers/Cryptography`.
- 4) Each DEK are then wrapped by using the KV, then serialized, and put into the columns metadata of the Parquet file.
- 5) The encrypted Parquet file is then written using Parquet Sharp.
- 6) Parquet file is then sent to the S3-compatible object store.
- 7) Columns, key IDs, and permissions are persisted in local PostgreSQL.

e) Decryption Flow:

Make into seq. diagram and make it more abstract and not code specific to give an overview of how OSWS handles it. Reference the diagram. Use pintora.

The decryption flow works as follows:

- 1) The client requests a Parquet file, via GET `/s3/{bucket}/{key}`.
- 2) OSWS fetches the encrypted Parquet file, first tries in local cache, then if not found, it goes to S3-compatible object store (see `OSWS.WebApi/Services/Services/S3ObjectFetcher`)
- 3) Wrapped DEKs are read from the Parquet column metadata.
- 4) For each wrapped DEK, OSWS checks the in-memory DEK cache. On a cache miss, it calls the KV decrypt method to unwrap the DEK and caches the result.
- 5) The user’s effective roles are computed via Listing 1.
- 6) The permitted columns are now decrypted, while the others are replaced by dummy columns⁴
- 7) The decrypted Parquet stream is returned to the client.

E. API

OSWS exposes three different endpoint groups.

a) S3-Compatible Endpoints:

integrate better with the rest of the text

A subset of S3 required endpoints is implemented to allow for the object store operations. These operations can be seen in Table III.

TABLE III
S3-COMPATIBLE ENDPOINTS PROVIDED BY OSWS

Method	Path	Description
GET	/	List buckets
GET	/ {bucket}	List objects in a bucket
PUT	/ {bucket}	Create bucket
DELETE	/ {bucket}	Delete bucket
POST	/ {bucket} ?delete	Delete multiple objects in bucket
HEAD	/ {bucket}	Retrieve bucket metadata
HEAD	/ {bucket} / {key}	Retrieve object metadata
GET	/ {bucket} / {*key}	Retrieve an object (with decryption and column filtering for Parquet)
PUT	/ {bucket} / {*key}	Store an object (with encryption for Parquet)
DELETE	/ {bucket} / {*key}	Delete an object

Currently, non-Parquet files pass through the encryption step, and the endpoints defined are only to allow operations made by most query engines, and thus suffice for making an MVP. In the future, this should be extended to also allow non-Parquet files to be encrypted and then include the KEK reference within the Parquet file referencing it, and the endpoints should also be extended.

b) Application Endpoints:

integrate better with the rest of the text

OIDC-protected endpoints for the web frontend. See the endpoints in Table IV.

TABLE IV
APPLICATION ENDPOINTS PROVIDED BY OSWS FOR WEB FRONTEND

Method	Path	Description
GET	/api/me	Current user profile (triggers JIT provisioning)
GET	/api/credentials	List the user's S3 credentials
POST	/api/credentials	Create a new S3 credential (access key + secret key)
DELETE	/api/credentials/ {id}	Revoke a credential (soft delete)

F. Limitations

OSWS set out to provide columnar access control to external query engines and “fully managed all-in-one cloud

platforms” by encrypting columns individually using PME, provided through the Parquet Sharp library. But due to early design decisions, this was not possible. As OSWS uses Parquet Sharp and stores cryptographic keys within the metadata of the Parquet files, a few issues appear with some query engines. First, Parquet Sharp uses Apache Arrow, which itself supports writing metadata.[27] This in itself makes sense, as it is possible by looking at e.g. `created_by` to see who created the Parquet file.[28] But as Parquet Sharp needs to write when encrypting the Parquet file, the file is essentially changed with new metadata, and it will no longer be the same Parquet file that the query engine inserted. This results in the fact that if the query engine stores metadata for ranges that are incorrect, which break their range requests. On top of that, OSWS saves cryptographic-related metadata within the column metadata, and that is yet another modification. As such, for this to work with all query engines, it would have to be ensured that all ranges of data remain the same, also after encryption.

The same issue is likely prevalent on most “fully managed all-in-one cloud platforms”, but it is not possible to test, due to it needing to have a URL whitelisted by the providers to test the OSWS system. But given documentation of e.g. Snowflake, it seems to point towards incompatibility.[29], [30]

In Section VII, this will be discussed in more depth and possible solutions to bridge these two design incompatibilities.

Lastly, initially, it was intended to make OSWS return the encrypted columns. But due to the limitations in Parquet Sharp, it was not possible, and the current solution is instead to create null/default (dummy) data entries for the column. There is an unimplemented method ready to encrypt this dummy data to simulate the idea of returning encrypted data, but it was not implemented due to being redundant, as OSWS already had the previously mentioned limitations.

V. METHODOLOGY

To ensure OSWS could live up to the intended use-cases and to check the viability of it, both benchmarking and e2e tests are quite important. This section will cover both the setup of the benchmark and the setup of the e2e test.

All source code is available at github.com/lucasfth/osws, and also describes how to get it running.

A. E2E Tests

As one of the goals of the OSWS was to enable third-party query engines to connect without any modification, a test confirming this behaviour is needed.

An e2e test will be created that asserts different query engines can connect and perform queries as normal, with OSWS providing column-level security based on the test user's roles. The e2e test suite will contain a seeding step that sets up the necessary users and roles. Python scripts will be used to orchestrate setting up roles and permissions, and asserting that the fetched data does not contain unauthorized data. The test will use the “Titanic” sample dataset⁵ and permissions will be configured so two roles have limited access, but to

⁴Due to limitations in Parquet Sharp, it is not possible to leave the encrypted column alone, and thus has to be replaced.

⁵<https://www.agentsfordata.com/app?sample-id=titanic&tab=local>

TABLE V
MAIN CONFIGURATIONS OF BENCHMARKS THAT ARE TO BE RUN

Category	Benchmark	Configuration	Measurements
E2E	S3 (Spaces)	No. rows: 1,000, 10,000, 250,000, 1,000,000 (50 columns)	Latency
E2E	OSWS +encryption +caching	No. rows: 1,000, 10,000, 250,000, 1,000,000 (50 columns)	Latency
E2E	OSWS +encryption -caching	No. rows: 1,000, 10,000, 250,000, 1,000,000 (50 columns)	Latency
E2E	OSWS -encryption -caching	No. rows: 1,000, 10,000, 250,000, 1,000,000 (50 columns)	Latency
Micro	Permission Service	No. flat hierarchy roles: 4, 64, 256	Latency
Micro	Permission Hierarchy	No. hierarchical depth roles: 0, 4, 16, 64	Latency
Micro	Unwrapping	Key size: 128,192, 256 bits	Latency
Micro	Decryption	No. rows: 1,000, 10,000, 250,000, 1,000,000, 2,000,000 (50 columns)	Latency

different columns. One role will inherit both and have direct access to all other columns, thus having full access. This will then test the role hierarchy feature. Then, the following will be used to fetch the Parquet file through different users and assert that only the permitted columns are viewable:

Python script: A Python script using an S3 client like `boto3` and a Parquet reader like `pyarrow` to fetch a file from OSWS.

DuckDB: A DuckDB instance will query the parquet file through OSWS.

Apache Spark: An Apache Spark instance (through PySpark) will query the parquet file through OSWS.

Each of the different tools should assert that only the permitted data is viewable.

B. Benchmarking

All benchmarking of OSWS was run on a Digital Ocean VM with access to 8 cores and 16GB of DRAM, see Appendix 1.b for full specification overview, together with Digital Ocean’s own Object Store (Spaces), which is S3-compatible. This helps ensure minimal latency between the compute and storage. For benchmarks using a key vault, Azure KV was used. All services were kept in the same region to further ensure minimal network latency. The version of OSWS used was [V2026.0.0 Alpha](#).

The test parquet files used in the following benchmarks can be defined as follows:

- *Parquet Size Tiny:*~500KB Parquet file with 50 columns and 1,000 rows.
- *Parquet Size Small:*~5MB Parquet file with 50 columns and 10,000 rows.
- *Parquet Size Medium:*~125MB Parquet file with 50 columns and 250,000 rows.
- *Parquet Size Large:*~500MB Parquet file with 50 columns and 1,000,000 rows.
- *Parquet Size X-Large:*~1GB Parquet file with 50 columns and 2,000,000 rows.

The five Parquet files are auto-generated by a script that is also available in the source code. The configuration of the files was chosen to reflect a somewhat realistic distribution of different sizes of Parquet files. Various openly available datasets were compared to settle on the choice of 50 columns. For example, several real-world example datasets are available

from [ClickHouse](#) contain much fewer than 50 columns.[31] Another is a dataset collecting 1.1 billion taxi and Uber rides in New York City, which contains 51 columns.[32] 50 columns seemed like a realistic number of columns for a large OLAP dataset, and was used for the sizes.

a) Micro Benchmarks:

The micro benchmarks that will be run will show how each major part of OSWS performs to get insights into how much they add of latency. The following micro-benchmarks will be run:

- 1) Permission Service
- 2) Permission Hierarchy
- 3) Unwrapping
- 4) Decryption

Permission Service will measure the latency of how long it takes to get a user’s effective roles to the given columns. The columns stay at a constant 100, the hierarchy (meaning roles pointing to roles) stays at zero, the roles assigned to the user are varied from four, 64, and 256 roles, and they are each assigned, on average, access to 80% of each of the columns.

Permission Hierarchy will measure the latency of how long it takes to get a user’s effective roles, but instead, roles are now linked to other roles and lastly to columns. The number of roles assigned is constantly at one, while the number of roles assigned in depth is varied from 0, 4, 16, and 64 and again, on average, each role is assigned access to 80% of the columns.

Unwrapping will measure the e2e latency for reading a tiny Parquet file (the small Parquet files are also run but ignored, but can be seen in Appendix 1.c) with a cold cache. The varying key sizes of 128, 192, and 256 bits will primarily affect the symmetric AES-GCM decryption.

Decryption will measure the latency for reading and decrypting a Parquet file, and the keys for decrypting will have been warmed and thus stored in cache; thus, KV is not taken into account here. The Parquet sizes that will be used are the same as with e2e benchmarking, being the tiny, small, medium, large, and extra large.

All of the micro-benchmarks were run 15 times to obtain statistically meaningful results.

See Table V for an overview of the benchmarks and their configurations, which will be run.

b) E2E Benchmarks:

The micro benchmarks show the individual parts' contribution to latency, but the overall performance of OSWS is more important to identify if, from a performance perspective, OSWS is usable within production environments.

As OSWS inherently implements different design choices compared to vanilla S3, no standardized benchmarking has been chosen, such as *Warp* which is otherwise used by multiple other Fortune 500 companies.[33], [34] The reasoning for this is that OSWS implements its own caching layer and uses a KV; as such, the benchmark would have to ensure that it tests both the throughput of N instances of OSWS running and that the configuration of caching and encryption is considered thoroughly.

OSWS also encrypts each column, resulting in the sizes of the Parquet files having a lot of significance for the end throughput, and Warp does not vary its Parquet file size. This results in inherently different throughputs and numbers, which are hard to compare OSWS against anything else.

Lastly, OSWS, as mentioned, does not support proper **range requests**, and it is then apparent that S3 will perform even better for ranged requests, where, regardless of the range or full file request, S3 will perform close to the same.⁶

Python scripts will then run Parquet puts and gets against the following setups:

- Directly against an S3-compatible store
- Against OSWS with encryption and with caching
- Against OSWS with encryption and no file cache
- Against OSWS with encryption and no DEK cache
- Against OSWS with no encryption and no caching

The metric for this will be latency, and for each file, the benchmark will be run for a total of 10 times, both cold and warm, allowing comparison of how slow or speed differences occur with and without caches.

VI. RESULTS

This section will cover the results of running both the benchmarks and e2e tests described in Section V, and describe their significance.

A. E2E Tests

An e2e test suite was written to test Python using PyArrow, DuckDB and PySpark against OSWS as an S3 endpoint. No modifications were made to the tools outside of changing the S3-endpoint to an instance of OSWS. The test suite shows that all three tools are successfully able to query a Parquet file and perform operations on the result. The users are identified through their credentials, and columns that the users' roles do not give access to are masked correctly.

This demonstrates that OSWS works as a drop-in replacement for S3 when working with query engines that fetch Parquet files directly and do not store metadata.

The same e2e test was tried to be done for DuckLake, but due to how OSWS modifies the Parquet file, it was not supported.

B. Benchmarking

a) Micro-Benchmarks:

In Figure 4, all box-plots increase in duration whenever their role depth, flat roles, or the number of rows increases (Note that the y axis are not necessarily using the same scale, and three of them are log-scaled). The latency for Permission Hierarchy Benchmark, see Figure 4c, and Permission Service Benchmark, see Figure 4d, both have a minor contribution to latency, which is expected for RBAC-related calls. These will not be referenced more as optimizations for these would not provide much benefit.

The major contributors to latency are decryption, see Figure 4a, and DEK unwrapping, see Figure 4b.

DEK unwrapping is close to constant. But the time is not something that can be improved much besides changing the KV service provider to a faster alternative, and then, as OSWS currently does, cache the unwrapped DEK heavily to use KV as little as possible. The values can be seen in more details in more details in Appendix 1.c, but essentially the $p95$ time for unwrapping key size 256 of a tiny file is 4.11 seconds, with a mean of 3.92 seconds, meaning that caching the unwrapped DEKs is essential, as if not cached this time would have to be added to each time a column has to be decrypted.

Decryption is the other main contributor. With the tiny file, it has a $p95$ of 17.1 milliseconds and a mean of 13.6ms, whilst for the medium and x-large, they are respectively, $p95$ 1.8 seconds, mean 1.68 seconds, and $p95$ 15.5 seconds, mean 14.7 seconds. So, essentially, to make OSWS usable for larger files, changes are needed in the decryption step.

The large numbers for the decryptions are due to a few design decisions. First of *V2026.0.0 Alpha* relies on Parquet Sharp version 21.0.0, which, as mentioned, makes it necessary, regardless of permission level, to copy over the whole Parquet file. To optimize this, a custom low-level reader would be needed, which could decrypt the columns in place. This would help, as the current decryption micro-benchmark also takes the copying into account, though the main improvement would be seen when columns are not authorized to the client, and can thus be left as is. So instead of the computations being directly related to the number of Parquet columns, they will be related to the number of authorized columns.

Decryption: $O(|\text{columns}|) \geq O(|\text{columns}_{\text{authorized}}|)$

b) E2E Benchmarks:

To begin with, comparison of cold GET latency with and without the DEK cache enabled is shown in Figure 5. It can be seen that a ceiling is hit at ~ 600 seconds for the large and x-large Parquet files, without DEK cache enabled. This can be attributed to the benchmarking timing out. The DEK cache can be seen to be invaluable on even small Parquet files, and on larger files, the system gets unusably slow without the DEK cache. At first, it seemed surprising that there was such a significant difference between a cold GET and a warm GET. For the tiny GET, it was close to 120 times and 3562 ms slower on average, and for the medium, it was close to 4 times and 4148ms slower.⁷ However, it makes sense given

⁶A solution to this is proposed within Section VII.

⁷Calculated from the results shown in Appendix 1.d




Fig. 4a. Decryption boxplot distribution per parameter (logarithmic y axis)




Fig. 4b. KeyUnwrap box-plot distribution per parameter set (tiny corpus only; linear y axis)




Fig. 4c. PermissionHierarchy box-plot distribution per parameter set (logarithmic y axis)




Fig. 4d. PermissionService box-plot distribution per parameter set (logarithmic y axis)

Figure 4. Box-plots for micro-benchmarks. Note that the y axes are not the same and some of the plots are logarithmic

that the DEK cache is warmed up as the file is being read, since reading a row chunk will populate the DEK cache with the given columns of that chunk. Since the whole column is encrypted with the same DEK, once that column is reached again in a new row chunk, the cache is warm. This shows that the DEK cache is a must-have for OSWS. For that reason, the “no DEK cache” configuration is omitted from here on.

Figure 6 shows the results of the rest of the e2e benchmark suite. The large and x-large Parquet file size results should be omitted in the evaluation, as the benchmarks became




Figure 5. DEK cache impact on median end-to-end latency (logarithmic y axis)




Fig. 6a. Median GET latency on warm cache




Fig. 6b. Median GET latency on cold cache




Fig. 6c. Median PUT latency

Figure 6. Plots showing median e2e latency measured during benchmarks. Grouped by file size. Note that legend is shared. (logarithmic y axis)

increasingly chaotic. This can be due to various factors, but is reasoned to be due to data fragmentation increasing over the benchmark runs. Lastly, reading from the data lake and the encrypted file cache might also have some unforeseen consequences, as they respectively use a memory stream and a file stream. But as the main latencies in the system were related to encryption/decryption and KV, this was down-prioritized.

Looking at Figure 6a for a warm GET, it is again clear that the DEK cache clearly decreases latency. On smaller files, all OSWS configurations come close to the direct Digital Ocean fetch, though the overhead is significant. Interestingly, the “no encrypt” is seemingly a bit *slower* than the other configurations for the “tiny” and “small” configuration, even though it should just be the pure network overhead. However, disabling encryption also disables the file cache, which might be why it is slower – due to it having to go to Digital Ocean every time.

In Figure 6b cold GET is shown. Direct fetch is omitted here due to there being no “cold” path – this is the same as Figure 6a. Here, the “no encryption” configuration is an order of magnitude faster for the “tiny” and “small” configurations, which highlights the overhead of the decryption step. However, as the files get larger, the network overhead starts to

show as well, when the “no encryption” configuration starts to approach the other configurations.

Generally, both Figure 6a and Figure 6b also show that the file cache is mostly useful when files are large; this is not surprising, as the encrypted file cache only helps with reducing S3-fetches, as the file must still be decrypted again. But when the network cost starts getting visible, the file cache matters.

Figure 6c shows the PUT latency. Unsurprisingly, there is massive overhead on the “tiny” configuration, as the direct upload is in millisecond-level, while OSWS must spend time encrypting the file and making calls to KV. Again, as files get larger, the relative overhead reduces as network transfer becomes a factor.

VII. DISCUSSION

OSWS set out to provide columnar access control to 3rd party query engines and “fully managed all-in-one cloud platforms” by encrypting columns individually using PME, provided through the Parquet Sharp library. During the encryption flow, Parquet Sharp rewrites the entire Parquet file: it decompresses columns, copies them row-group by row-group into a new file, and adds cryptographic metadata such




Figure 7. Sequence Diagram showing the DuckLake range request error case. File size and byte ranges are for illustrative purposes only.

as wrapped DEKs and KEK references, as well as modifying existing metadata, e.g. the `created_by` field. The resulting file is larger than the original and no longer byte-compatible with it. This single root cause is responsible for the majority of the incompatibilities identified during the final iteration of benchmarking and end-to-end testing, and this section addresses those issues, defines what should have been done instead, and outlines the performance improvements available to a future OSWS system.

A. Limitations of the Current Design

The core issue introduced by Parquet Sharp is that the written Parquet file no longer matches the file seen by clients. Three concrete effects follow from this:

- **Modified metadata:** Cryptographic metadata – wrapped DEKs, KEK references, and modified fields such as `created_by` – is appended to the file footer and column metadata, increasing the total file size.
- **Decompressed columns:** Parquet Sharp requires columns to be decompressed before they can be copied. **Not entirely correct, we just use default settings, which is uncompressed. It could have been compressed.** The re-serialized file is therefore stored without compression, further inflating the file size relative to the original.

- **Sequential, full-file processing:** Parquet Sharp reads the entire file regardless of the byte range a client requests. Keys are unwrapped one by one after the full file has been fetched, and columns are then copied sequentially. Each step waits for the previous to complete.

Together, these mean that OSWS cannot return the original encrypted bytes for non-authorized columns, cannot serve range requests correctly once file size has changed, and cannot overlap any part of its fetch, decrypt, and response pipeline.

B. Impact on Query Engines

The effect of the modified file size differs depending on whether a query engine is *schema-on-read* or *schema-on-write*.

Schema-on-read engines such as DuckDB and Spark are unaffected by the size change. When they fetch a file, they first issue a HEAD request to obtain the current file size, and then perform a range request for the footer using that size. Since the HEAD request returns the new encrypted file size, the subsequent range request is correct.

Schema-on-write tools that maintain a catalogue database, such as DuckLake[35], exhibit broken behaviour. When DuckLake writes a Parquet file, it records both the file size and the footer byte range in its catalogue, allowing it to skip the HEAD request on subsequent reads and fetch the footer

directly. For OSWS, however, the written file changes after the client PUT completes, so the recorded ranges are no longer valid. See Figure 7 for an illustration of this issue. The following describes it in more detail.

To begin, DuckLake uses its internal Parquet writer “duck-lake-writer” to create and PUT a Parquet file of 1000 bytes. Then, OSWS uses its writer, ParquetSharp, to create an encrypted copy. This is now a whole new Parquet file with encrypted columns, new metadata and possibly different encoding due to the different writers - illustrated by a file size of 2000 bytes. This file is then stored in the storage backend.

When DuckLake later queries the file, it issues a range request of bytes 500–1000, being the range it thinks contains the footer.

Since OSWS must fetch the entire file to be able to decrypt, the 2000 byte file is first fetched, and then decrypted. However, on decryption, the file is again copied into a whole new Parquet file. Now, the values are the same as the original, DuckLake-written 1000 byte file, but again, since the writers are not the same, the metadata and encoding might be different. This is illustrated with a file size of 1250 bytes.

To serve the range request, OSWS then slices *this* file in the range 500-1000, but given the different file size, this results in a wrong range. This is then sent to DuckLake, which throws an error due to not finding the expected data.

This eliminates DuckLake and probably Snowflake from being compatible with OSWS in its current form. The root cause is not the specific metadata that is added, but the fundamental incompatibility between recording file ranges at write time and then rewriting the file.

C. Proposed Redesign

The issues described above share a common fix: OSWS should neither store cryptographic metadata inside the Parquet file, nor rewrite the file at all. This requires two coordinated changes: an internal metadata store, and a custom Parquet reader and writer using AES CTR.

a) Internal Metadata Store:

Cryptographic metadata should be stored in an internal SQL database within OSWS rather than in the Parquet file column metadata. For each Parquet file, the database should contain:

- 1) Internal file identifier for the footer⁸
- 2) Byte range of the start and end index
- 3) Internal file identifier for the column index⁸
- 4) Wrapped DEK
- 5) KEK reference
- 6) Row chunk size

When a read request arrives, OSWS can retrieve the file from the object store while concurrently querying the local database for matching entries under 1. It can then cross-reference the client’s RBAC permissions and match authorized columns against 3, fetching wrapped DEKs from KV using 5. For range requests, OSWS can use 2 to determine which columns the requested byte range covers, fetching only those

from S3 and requesting only the relevant DEKs — rather than fetching the entire file.

b) Custom Reader and Writer with AES CTR:

The current dependency on Parquet Sharp should be replaced with a custom Parquet reader and writer that encrypts and decrypts in place using AES CTR. AES CTR uses no padding, which means the encrypted form of a column chunk is the same size as the plaintext, preserving all byte offsets and leaving the file size unchanged. Since column metadata is no longer embedded in the file, compression can also be left intact, as there is no longer a need to decompress columns before encryption.

This enables three concrete improvements over the current design:

- 1) *Asynchronous key pre-fetching*: While the encrypted Parquet file is being streamed from the object store, OSWS can concurrently request DEK unwrapping from KV. By the time the file has arrived, the unwrapped keys are ready, removing KV latency from the critical path.
- 2) *In-place decryption*: Rather than copying the entire file, only the requested column chunks are decrypted, using AES CTR’s preserved byte offsets to locate them directly.
- 3) *True streaming*: As soon as a row chunk arrives and its key is ready, decryption can begin immediately, without waiting for the entire file to be fetched first.

AES CTR does carry known considerations, discussed by Helger et al.[37]:

- No integrity: CTR provides no message integrity, but this can be handled by a MAC layer.
- Error propagation: Bit flips are localized and do not propagate. This should be addressed at a separate layer.
- Stateful encryption: Keys must not be reused, which is already enforced by OSWS generating a new DEK per column.
- Counter reuse: Counter values must not be reused; given per-column key generation, this is satisfied.

Additionally, knowing the column chunk size and row chunk size from the internal store allows OSWS to identify the precise offset within a range query, fetch only the relevant bytes from S3, and decrypt only that portion, which substantially reduces the overhead compared to the current full-file fetch.

c) Parallelization:

With in-place encryption and decryption, row group processing becomes an embarrassingly parallelisable operation. OSWS can spawn threads to process independent row groups concurrently rather than sequentially, which should reduce processing overhead substantially. This was not feasible with Parquet Sharp, which provides no API for writing bytes at a specific offset.

D. Encrypted Parquet Storage Cache

When Parquet files are fetched from S3, they are stored in local storage in their encrypted format. As seen in Section VI.B.b, this yields negligible improvement for smaller files, with more significant gains on larger ones. With proper

⁸Defined within [36]

range request support, this improvement is likely to diminish further, since most requests will cover only a portion of a file rather than fetching it in full. Caching in unencrypted form in DRAM would offer additional performance gains, but it introduces complexity around column-level access control and memory pressure. This is therefore considered a later-stage optimization to be evaluated on a future OSWS system.

E. Correct Decisions & OSWS Future

Despite the design issues outlined above, several decisions made in the current OSWS implementation are considered sound and should be carried over to a V2 system.

a) Envelope Encryption:

Envelope encryption should be retained. While it carries a significant initial overhead, this is substantially mitigated by the decrypted DEK cache. It also enables straightforward key rotation, which is an important property for future security improvements. A fully fledged key vault inside OSWS would be an alternative that eliminates the KV dependency, but would require significant additional security measures.

b) DEK Cache:

The decrypted DEK cache must be carried over to any future version that uses a KV store. As seen in both Section VI.B.b and Section VI.B.a, DEK unwrapping carries a significant overhead, and the cache is the primary mechanism for keeping it off the critical path on repeated accesses.

c) RBAC Metadata Storage:

The microbenchmarks in Section VI.B.a showed that authorizing users via the RBAC database contributes negligible overhead relative to cryptographic operations and network transfer. A PostgreSQL database with the implemented schema is therefore a sound choice for RBAC metadata storage going forward.

d) Role Management:

Role management — creating and assigning roles, granting permissions — using the implemented query editor provided an intuitive and familiar interface over the admin API. While this is a wrapper that could be implemented in many ways, it is a valid option for future versions.

e) Known Security Considerations:

Two implementation-level security issues are acknowledged in the current OSWS release, documented in [KNOWN_ISSUES.md](#).

First, S3 credential secret keys are stored in plaintext in the PostgreSQL database, as the SigV4 authentication handler requires the raw key for HMAC derivation. If the database is compromised, all S3 credentials become immediately usable. A mitigation would be to encrypt secret keys at rest using Azure Key Vault, decrypting on read during SigV4 verification.

Second, the RBAC admin flag is provisioned just-in-time from OIDC claims on every login. If the OIDC provider exposes this as a self-service field, users could grant themselves admin access. This is dependent on the OIDC provider configuration and can be mitigated by using a dedicated admin identity provider or restricting claim sources.

VIII. CONCLUSION

The PoC OSWS demonstrates the possibility of adding fine-grained role-based access control in data lakes by using encryption with native Parquet files, and thus enforcing it within the Object Store. Query engines that do not cache Parquet sizing metadata, such as DuckDB, can use OSWS without modifications by pointing their S3 URL to OSWS.

This thesis had the following five objectives:

- 1) Create OSWS PoC using an underlying S3-compatible Object Store
- 2) Ensure OSWS enforces access control using RBAC outward using encryption
- 3) Make it compatible with “fully managed all-in-one cloud platforms”
- 4) Make it compatible with query engines
- 5) Benchmark OSWS to measure the latency it adds.
- 6) Reflect on design decisions and which alternatives would have been better

Only four of the objectives were partially completed, as 3 were not possible due to limitations of whitelisting. Fine-grained RBAC with column-level encryption was achieved in OSWS, proved by an e2e test, and addressing objectives 1 and 2. 3rd party query engines not using cached metadata, such as DuckDB, could have OSWS dropped in as an S3 replacement without modifying them, partially solving objective 4.

The whole system was benchmarked, and showed that a DEK cache is non-negotiable, even with small Parquet files; the latency improved dramatically when using it. This solved objective 5. But generally, it showed that OSWS in its current state is not production-ready, due to the design choices made.

The results of the current state of OSWS are better seen as design proof that such a system could work. The issues identified were a lot of sequential work, modified metadata, and unsupported range requests internally. But these issues have been reflected upon and potential solutions proposed within Section VII, thus addressing objective 6.

I. APPENDIX

A. Encrypted Parquet Footer




Fig. 1. Illustration of the Parquet file structure using an Encrypted Footer. From [18].

B. VM Specification

```
root@ubuntu-s-1vcpu-1gb-35gb-intel-fral:~# inxi -Fxz
System:
  Kernel: 6.8.0-111-generic arch: x86_64 bits: 64 compiler: gcc v: 13.3.0
  Console: pty pts/6 Distro: Ubuntu 24.04.4 LTS (Noble Numbat)
Machine:
  Type: Kvm Mobo: DigitalOcean model: Droplet v: 20171212 serial: <filter> BIOS:
  DigitalOcean
  v: 20171212 date: 12/12/2017
CPU:
  Info: 8-core model: DO-Premium-AMD bits: 64 type: MCP arch: Zen 2 rev: 0 cache: L1:
  512 KiB
  L2: 4 MiB L3: 16 MiB
  Speed (MHz): avg: 1996 min/max: N/A cores: 1: 1996 2: 1996 3: 1996 4: 1996 5: 1996
  6: 1996
  7: 1996 8: 1996 bogomips: 31940
  Flags: avx avx2 ht lm nx pae sse sse2 sse3 sse4_1 sse4_2 sse4a ssse3 svm
Graphics:
  Device-1: Red Hat Virtio 1.0 GPU driver: virtio-pci v: 1 bus-ID: 00:02.0
  Display: server: No display server data found. Headless machine? tty: 110x29
  resolution: 1024x768
  API: EGL v: 1.5 drivers: kms_swrast,swrast platforms: active:
  gbm,surfaceless,device
  inactive: wayland,x11
  API: OpenGL v: 4.5 vendor: mesa v: 25.2.8-0ubuntu0.24.04.1 note: console (EGL
  sourced)
  renderer: llvmpipe (LLVM 20.1.2 256 bits)
Audio:
  Message: No device data found.
  API: ALSA v: k6.8.0-111-generic status: inactive
Network:
  Device-1: Intel 82371AB/EB/MB PIIX4 ACPI vendor: Red Hat Qemu virtual machine
  type: network bridge driver: N/A port: N/A bus-ID: 00:01.3
  Device-2: Red Hat Virtio network driver: virtio-pci v: 1 port: clc0 bus-ID: 00:03.0
  IF: eth0 state: up speed: -1 duplex: unknown mac: <filter>
  Device-3: Red Hat Virtio network driver: virtio-pci v: 1 port: clc0 bus-ID: 00:04.0
  IF: eth1 state: up speed: -1 duplex: unknown mac: <filter>
  IF-ID-1: br-cf3f443f079d state: down mac: <filter>
  IF-ID-2: docker0 state: down mac: <filter>
Drives:
  Local Storage: total: 320 GiB used: 19.4 GiB (6.1%)
  ID-1: /dev/vda model: N/A size: 320 GiB
  ID-2: /dev/vdb model: N/A size: 488 KiB
Partition:
  ID-1: / size: 308.93 GiB used: 19.28 GiB (6.2%) fs: ext4 dev: /dev/vda1
  ID-2: /boot size: 880.4 MiB used: 113.9 MiB (12.9%) fs: ext4 dev: /dev/vda16
  ID-3: /boot/efi size: 104.3 MiB used: 6.1 MiB (5.9%) fs: vfat dev: /dev/vda15
Swap:
  Alert: No swap data was found.
Sensors:
  Src: lm-sensors+/sys Message: No sensor data found using /sys/class/hwmon or lm-
  sensors.
Info:
  Memory: total: 16 GiB available: 15.62 GiB used: 1.03 GiB (6.6%)
  Processes: 180 Uptime: 4d 1h 2m Init: systemd target: graphical (5)
  Packages: 813 Compilers: N/A Shell: Bash v: 5.2.21 inxi: 3.3.34
```

Listing 3. VM specification used for the benchmarks of OSWS

C. Micro Benchmark Results

Benchmark: Decryption								
parameters	n	mean	stddev	min	p50	p95	p99	max
size=large	15	7.48s	313.4ms	7.12s	7.43s	8.17s	8.17s	8.17s
size=medium	15	1.68s	64.1ms	1.58s	1.67s	1.80s	1.80s	1.80s
size=small	15	65.7ms	4.8ms	58.7ms	64.2ms	75.0ms	75.0ms	75.0ms
size=tiny	15	13.6ms	1.2ms	12.0ms	13.4ms	17.1ms	17.1ms	17.1ms
size=xlarge	15	14.7s	452.1ms	14.1s	14.6s	15.5s	15.5s	15.5s
Benchmark: KeyUnwrap								
parameters	n	mean	stddev	min	p50	p95	p99	max
size=small,dek_bits=128	15	4.15s	166.5ms	3.94s	4.15s	4.55s	4.55s	4.55s
size=small,dek_bits=192	15	4.05s	198.0ms	3.76s	3.99s	4.37s	4.37s	4.37s
size=small,dek_bits=256	15	4.05s	186.2ms	3.73s	4.04s	4.53s	4.53s	4.53s
size=tiny,dek_bits=128	15	3.15s	152.0ms	2.86s	3.12s	3.44s	3.44s	3.44s
size=tiny,dek_bits=192	15	3.24s	214.5ms	2.96s	3.22s	3.76s	3.76s	3.76s
size=tiny,dek_bits=256	15	3.92s	113.8ms	3.67s	3.97s	4.11s	4.11s	4.11s
Benchmark: PermissionHierarchy								
parameters	n	mean	stddev	min	p50	p95	p99	max
hierarchy_depth=0	15	4.3ms	0.7ms	3.5ms	4.2ms	5.8ms	5.8ms	5.8ms
hierarchy_depth=16	15	6.8ms	2.3ms	5.1ms	6.2ms	14.8ms	14.8ms	14.8ms
hierarchy_depth=4	15	4.5ms	0.8ms	3.6ms	4.4ms	7.0ms	7.0ms	7.0ms
hierarchy_depth=64	15	18.5ms	2.0ms	14.7ms	18.8ms	23.6ms	23.6ms	23.6ms
Benchmark: PermissionService								
parameters	n	mean	stddev	min	p50	p95	p99	max
role_count=256	15	24.7ms	2.6ms	21.2ms	24.3ms	28.1ms	28.1ms	28.1ms
role_count=4	15	6.7ms	4.8ms	4.3ms	5.3ms	24.6ms	24.6ms	24.6ms
role_count=64	15	12.0ms	1.4ms	9.6ms	12.2ms	14.8ms	14.8ms	14.8ms

Listing 4. Textual representation for the results of running the micro benchmarks

D. E2E Latency Benchmark Results

PUT cache=n/a						
config	metric	tiny	small	medium	large	xlarge
s3-direct	mean	27ms	69ms	575ms	2031ms	4093ms
	stddev	11ms	24ms	112ms	39ms	114ms
	p50	24ms	60ms	535ms	2046ms	4049ms
	p95	57ms	140ms	846ms	2076ms	4317ms
	n	10	10	10	10	10
osws-encrypt-cache	mean	2208ms	2024ms	5024ms	14.2s	26.5s
	stddev	660ms	94ms	161ms	354ms	466ms
	p50	1994ms	2034ms	5046ms	14.3s	26.5s
	p95	4168ms	2236ms	5234ms	14.7s	27.2s
	n	10	10	10	10	10
osws-encrypt-no-file-cache	mean	2150ms	2090ms	5228ms	14.4s	26.1s
	stddev	584ms	81ms	139ms	353ms	557ms
	p50	1965ms	2081ms	5236ms	14.2s	26.2s
	p95	3884ms	2222ms	5495ms	15.2s	27.2s
	n	10	10	10	10	10
osws-encrypt-no-dek-cache	mean	2558ms	2327ms	5152ms	14.2s	25.3s
	stddev	567ms	264ms	207ms	197ms	772ms
	p50	2383ms	2263ms	5054ms	14.2s	25.4s
	p95	4246ms	2845ms	5554ms	14.6s	26.7s
	n	10	10	10	10	10
osws-no-encrypt	mean	172ms	153ms	1940ms	7456ms	14.4s
	stddev	339ms	9ms	104ms	291ms	244ms
	p50	48ms	151ms	1928ms	7476ms	14.4s
	p95	1186ms	172ms	2165ms	7868ms	14.9s
	n	10	10	10	10	10

Listing 5. Textual representation of E2E Latency Benchmark results for PUT

GET cache=cold

config	metric	tiny	small	medium	large	xlarge
osws-encrypt-cache	mean	3592ms	3708ms	6018ms	13.1s	29.4s
	stddev	193ms	94ms	131ms	391ms	11.0s
	p50	3558ms	3703ms	6024ms	13.2s	22.6s
	p95	3997ms	3893ms	6239ms	13.8s	46.9s
	n	10	10	10	10	10
osws-encrypt-no-file-cache	mean	3662ms	3819ms	6436ms	13.6s	29.8s
	stddev	116ms	160ms	283ms	248ms	11.3s
	p50	3648ms	3819ms	6362ms	13.6s	22.7s
	p95	3863ms	4166ms	7007ms	14.1s	47.2s
	n	10	10	10	10	10
osws-encrypt-no-dek-cache	mean	20.1s	16.2s	79.6s	600.6s	600.7s
	stddev	797ms	811ms	3156ms	369ms	387ms
	p50	20.1s	15.9s	79.0s	600.5s	600.9s
	p95	21.8s	18.3s	85.6s	601.2s	601.1s
	n	10	10	10	10	10
osws-no-encrypt	mean	45ms	151ms	1303ms	5017ms	10.5s
	stddev	17ms	17ms	50ms	107ms	491ms
	p50	38ms	153ms	1300ms	5016ms	10.2s
	p95	92ms	177ms	1390ms	5220ms	11.6s
	n	10	10	10	10	10

GET cache=warm

config	metric	tiny	small	medium	large	xlarge
s3-direct	mean	12ms	33ms	477ms	1884ms	3715ms
	stddev	2ms	3ms	44ms	104ms	162ms
	p50	11ms	33ms	490ms	1915ms	3737ms
	p95	16ms	41ms	555ms	2015ms	3987ms
	n	10	10	10	10	10
osws-encrypt-cache	mean	30ms	82ms	1600ms	6040ms	65.6s
	stddev	3ms	6ms	101ms	248ms	5962ms
	p50	29ms	82ms	1617ms	6097ms	62.9s
	p95	37ms	92ms	1769ms	6362ms	75.0s
	n	10	10	10	10	10
osws-encrypt-no-file-cache	mean	35ms	130ms	2298ms	9090ms	82.2s
	stddev	3ms	9ms	57ms	274ms	2451ms
	p50	34ms	129ms	2282ms	9025ms	81.9s
	p95	42ms	151ms	2419ms	9610ms	87.8s
	n	10	10	10	10	10
osws-encrypt-no-dek-cache	mean	16.0s	15.4s	78.1s	508.5s	600.4s
	stddev	2087ms	531ms	4794ms	140.2s	291ms
	p50	14.6s	15.4s	76.5s	600.3s	600.4s
	p95	20.3s	16.3s	91.6s	600.8s	601.0s
	n	10	10	10	10	10
osws-no-encrypt	mean	38ms	138ms	1303ms	5214ms	9953ms
	stddev	5ms	16ms	43ms	297ms	434ms
	p50	36ms	133ms	1302ms	5103ms	10.0s
	p95	51ms	167ms	1378ms	5742ms	10.6s
	n	10	10	10	10	10

Listing 6. Textual representation of E2E Latency Benchmark results for GET with warm and cold cache

REFERENCES

- [1] A. S. H. Trøstrup and L. F. T. Hanson, "Reviewing options for fine-grained Role-Based Access Control in Data Lakes." [Online]. Available: <https://github.com/lucasfth/reviewing-options-for-fine-grained-role-based-access-control-in-data-lakes>
- [2] S. Kumar, S. Yagati, C. Power, D. E. Culler, and R. A. Popa, "Membrane: A Cryptographic Access Control System for Data Lakes," *ArXiv*, 2025, [Online]. Available: <https://api.semanticscholar.org/CorpusID:281243750>
- [3] X. Shang *et al.*, "One Stone, Three Birds: Finer-Grained Encryption with Apache Parquet @ Large Scale," in *2022 IEEE International Conference on Big Data (Big Data)*, Osaka, Japan: IEEE, Dec. 2022, pp. 5802–5811. doi: [10.1109/BigData55660.2022.10020987](https://doi.org/10.1109/BigData55660.2022.10020987).
- [4] A. Jonker and M. M. Kosinski, "What is a data lake?." [Online]. Available: <https://www.ibm.com/think/topics/data-lake>
- [5] Deepstrike, [Online]. Available: <https://deepstrike.io/blog/what-is-jit-provisioning>
- [6] D. F. Ferraiolo and R. Kuhn, "Role-Based Access Controls," in *Proceedings of the 15th National Computer Security Conference (NCSC)*, Baltimore, MD, USA, Oct. 1992, pp. 554–563.
- [7] D. F. Ferraiolo, R. Sandhu, S. Gavrila, D. R. Kuhn, and R. Chandramouli, "Proposed NIST standard for role-based access control," *ACM Trans. Inf. Syst. Secur.*, vol. 4, no. 3, pp. 224–274, Aug. 2001, doi: [10.1145/501978.501980](https://doi.org/10.1145/501978.501980).
- [8] O. Foundation, "What is OpenID Connect." [Online]. Available: <https://openid.net/developers/how-connect-works/>
- [9] Microsoft, "What is Microsoft Entra?." [Online]. Available: <https://learn.microsoft.com/da-dk/entra/fundamentals/what-is-entra>
- [10] PocketID, "Introduction to PocketID." [Online]. Available: <https://pocket-id.org/docs/introduction>
- [11] AWS, "AWS Key Management Service." [Online]. Available: <https://docs.aws.amazon.com/kms/latest/developerguide/overview.html>
- [12] Azure, "Azure Key Vault Overview - Azure Key Vault." [Online]. Available: <https://docs.azure.cn/en-us/key-vault/general/overview>
- [13] AWS, [Online]. Available: <https://docs.aws.amazon.com/AmazonS3/latest/userguide/Welcome.html>
- [14] A. S. Foundation, "Apache Spark." [Online]. Available: <https://spark.apache.org/>
- [15] M. A. Zaharia, A. Ghodsi, R. Xin, and M. Armbrust, "Lakehouse: A New Generation of Open Platforms that Unify Data Warehousing and Advanced Analytics," in *Conference on Innovative Data Systems Research*, 2021. [Online]. Available: <https://api.semanticscholar.org/CorpusID:229576171>
- [16] Apache, "Parquet File Format." [Online]. Available: <https://parquet.apache.org/docs/file-format/>
- [17] A. S. Foundation, "Apache Thrift." [Online]. Available: <https://thrift.apache.org/>
- [18] Apache, [Online]. Available: <https://parquet.apache.org/docs/file-format/data-pages/encryption/>
- [19] Google, [Online]. Available: <https://docs.cloud.google.com/kms/docs/envelope-encryption>
- [20] Redis, "TTL." [Online]. Available: <https://redis.io/docs/latest/commands/ttl/>
- [21] E. Barker, *Recommendation for Key Management Part 1: General*, no. NIST SP 800-57pt1r4. 2016, p. NIST SP 800-57pt1r4. doi: [10.6028/NIST.SP.800-57pt1r4](https://doi.org/10.6028/NIST.SP.800-57pt1r4).
- [22] R. Anderson and T. Dykstra, "Tutorial: Create a Minimal API with ASP.NET Core." [Online]. Available: <https://learn.microsoft.com/en-us/aspnet/core/tutorials/min-web-api?view=aspnetcore-10.0>
- [23] msmbaldwin, "Azure data encryption at rest - Azure Security." [Online]. Available: <https://learn.microsoft.com/en-us/azure/security/fundamentals/encryption-atrest>
- [24] "Authenticating Requests (AWS Signature Version 4)." [Online]. Available: <https://docs.aws.amazon.com/AmazonS3/latest/API/sig-v4-authenticating-requests.html>
- [25] [Online]. Available: <https://bitwarden.com/resources/passkeys-are-phishing-resistant-mfa/>
- [26] peggyjs, "peggyjs/peggy – Parser Generator for JavaScript." [Online]. Available: <https://github.com/peggyjs/peggy>
- [27] [Online]. Available: <https://github.com/G-Research/ParquetSharp>
- [28] "Reading and writing Parquet files." [Online]. Available: <https://arrow.apache.org/docs/cpp/parquet.html>
- [29] Snowflake, "Introduction to external tables." [Online]. Available: https://docs.snowflake.com/en/user-guide/tables-external-intro?utm_source=DV-Blog%3Fsource&utm_cta=website-workload-ai-ml-timely-content-building-a-chatbot-app-with-rag-and-feature-store
- [30] Snowflake, "External table files." [Online]. Available: https://docs.snowflake.com/en/sql-reference/functions/external_table_files
- [31] ClickHouse, "ClickHouse Example Datasets." [Online]. Available: <https://clickhouse.com/docs/getting-started/example-datasets>
- [32] T. W. Schneider, "Analyzing 1.1 Billion NYC Taxi and Uber Trips, with a Vengeance." [Online]. Available: <https://toddwtschneider.com/posts/analyzing-1-1-billion-nyc-taxi-and-uber-trips-with-a-vengeance/>
- [33] "Companies using MinIO in 2026 | Landbase." Accessed: Mar. 23, 2026. [Online]. Available: <https://data.landbase.com/technology/minio/>
- [34] "Companies that use Minio (2,450)." Accessed: Mar. 23, 2026. [Online]. Available: <https://theirstack.com/en/technology/minio>
- [35] DuckDB Foundation, "ducklake_data_file." [Online]. Available: https://ducklake.select/docs/stable/specification/tables/ducklake_data_file.html
- [36] "Parquet Modular Encryption." [Online]. Available: <https://parquet.apache.org/docs/file-format/data-pages/encryption/>
- [37] H. Lipmaa, P. Rogaway, and D. Wagner, "Comments to NIST concerning AES-modes of operations: CTR-mode encryption," p. , 2000.